

ROHIT ANANTHAN

Data Scientist — Experimentation & Causal Measurement

rohitananthan@live.com · (202) 989-9596 · rohitananthan.info · linkedin.com/in/rohit-ananthan · github.com/Ramidoz

SUMMARY

Data scientist with 4+ years across experimentation, causal measurement, and ML-driven product and campaign analytics. Designs A/B tests with primary and guardrail metric frameworks, applies quasi-experimental cohort methods when a clean test isn't feasible, and owns the SQL/Python pipelines, dashboards, and automation that turn one-off analyses into reliable, self-serve tools. Deep background in lifecycle and campaign analytics: segmentation, churn, reactivation, and incrementality measurement.

EXPERIENCE

AI Engineer — Vdart Inc. | Remote | Jun 2026 – Present

- Building a multi-agent system that automates the full RFP lifecycle for an enterprise healthcare client — intake, routing and assignment, QA review, and rules-based scoring against current evaluation criteria — converting a recurring manual workflow into a repeatable system targeting ~85% less manual processing time.
- Engineered retrieval over an Azure AI Search vector index with MongoDB persistence to ground scoring in the latest rules; accelerated delivery across a React/Node.js stack using GitHub Copilot-orchestrated coding agents.

Data Scientist — Invision Global Tech Inc. | Remote | Feb 2026 – Jun 2026

- Built evaluation harnesses for production LLM workflows, including reusable prompt patterns and tool-calling templates that compressed turnaround on recurring engineering work.
- Prototyped ML and analytics solutions for client engagements using Python, FastAPI, SQL, and Claude Code-assisted development.

Data Scientist — Community Dreams Foundation | Remote | Feb 2025 – Feb 2026

- Built and deployed a production voice FAQ agent (Gemini Pro, LangChain, LlamaIndex) and defined its evaluation framework — answer accuracy, task completion, deflection rate, and latency — so every pipeline change was judged against a fixed test set instead of subjective review.
- Ran three controlled evaluation cycles comparing RAG and model-pipeline iterations, improving answer accuracy ~22%; selected winning configurations on guardrails (latency, task completion, failure behavior) rather than raw accuracy alone.
- Designed and shipped scalable ML pipelines on GCP Vertex AI with Dataflow, Airflow, and BigQuery, cutting latency on recurring data flows by 60% and standing up MLOps monitoring so stakeholders could act on recurring reads with confidence.

Data Analyst — Kameleon Technologies | India | May 2021 – Jul 2023

- Designed and analyzed A/B tests on campaign and product-experience rollouts for a B2B retail platform, defining primary metrics (engagement, reactivation) alongside guardrail metrics to protect the broader customer experience; experimentation contributed to 12–15% gains in engagement and reactivation.
- Applied difference-in-differences-style cohort analysis to measure incremental campaign impact — comparing pre/post outcomes for exposed cohorts against comparable unexposed cohorts — separating true campaign lift from baseline client differences and informing which campaigns to continue, modify, or cut.
- Shipped customer segmentation (K-means, RFM) and churn models in Python that targeted retention and lifecycle campaigns, improving 90-day retention by ~12–15% and guiding marketing-investment decisions.
- Owned ETL pipelines and analytics infrastructure processing 5M+ daily transactions, replacing manual reporting as the source of truth across product, revenue, and fraud (Neo4j graph analytics, XGBoost) workflows.

SKILLS

- **Experimentation & Causal Measurement:** A/B test design, primary/guardrail metric frameworks, difference-in-differences-style cohort analysis, incrementality and pre/post measurement, evaluation harness design, statistical analysis
- **Data Science & ML:** Python, SQL, scikit-learn, XGBoost, segmentation (K-means, RFM), churn and fraud modeling, MLflow, PySpark
- **Pipelines, Dashboards & Automation:** Airflow, dbt, BigQuery, Vertex AI, Dataflow, Snowflake, MongoDB, Neo4j, FastAPI, self-serve dashboards and reporting automation
- **GenAI & Productivity Tooling:** RAG, LLM evaluation, LangChain, GitHub Copilot, Claude Code, Cursor; AWS Certified AI Practitioner

EDUCATION

- **MS, Information Systems** — University of Maryland, Robert H. Smith School of Business | Aug 2023 – Dec 2024
- **BTech, Engineering** — Indian Institute of Information Technology | 2018 – 2022 | ML paper at ICRAM 2022, IIT Jodhpur (Springer)